



GOVERNANCE INTELLIGENCE REPORT

Trust Evaluation — PDF Test

Deep Research Analysis · VeldrixAI

Report Name	Cobalt Nexus
Report ID	VX-TEST-0001
Generated	May 23, 2026 17:08 UTC
Evaluated Model	llama-3.1-70b-instruct (NVIDIA NIM)
Evaluation Window	Last 30 Days
Total Evaluations	0
Tenant / Org	VeldrixAI Platform
Pillar Version	v2.1.0

73.0

/ 100 · CAUTION

Overall Trust Score

Weighted average across five VeldrixAI trust pillars: Safety (25%), Hallucination (25%), Bias (20%), Prompt Security (15%), Compliance (15%).

This report was automatically generated by the VeldrixAI Runtime Trust Engine. It contains AI-evaluated risk analysis, policy compliance findings, and enforcement recommendations based on live production traffic. All scores are computed deterministically using the VeldrixAI five-pillar evaluation framework. This document is confidential and intended solely for the authorized recipient.



1. EXECUTIVE SUMMARY

Test evaluation. Hallucination and Compliance scored below target.

73.0	1	0	0.0%
Overall Trust Score	Total Evaluations	Interventions Made	Pass-Through Rate

2. TRUST PILLAR ANALYSIS




Figure 1: Trust Pillar Radar (left) and Score Comparison (right). Dashed lines indicate minimum threshold (70) and target performance (90).

Pillar Score Detail

Pillar	Weight	Score	Status	Contribution
Safety	25%	91.0	✓ PASS	22.75
Hallucination	25%	62.0	✗ FAIL	15.50
Bias	20%	78.0	■ CAUTION	15.60
Prompt Security	15%	88.0	✓ PASS	13.20
Compliance	15%	55.0	✗ FAIL	8.25



3. ML SIGNAL ANALYTICS

The following charts characterise the model's safety profile at inference time. The scatter plot positions each trust pillar in weight–score space; bubble area encodes the weighted contribution to the overall trust score. The verdict pie summarises per-pillar classification outcomes. The risk matrix heatmap provides a columnar breakdown of score, weight, contribution, and inverted risk index for rapid triage.




Figure 2: Pillar risk scatter (left) — bubble area ∝ weighted contribution; shading bands indicate Fail (<70), Caution (70–84), and Pass (≥85) zones. Pillar verdict distribution (right) — proportion of pillars in each risk category.




Figure 3: Risk matrix heatmap. Columns: Score (0–100), Weight %, Weighted Contribution, Risk Index (0 = trusted, 100 = critical). Colour gradient: green → amber → red. Risk Index column uses an inverted gradient — darker red indicates higher risk.



4. FINDINGS & RISK ASSESSMENT

Pillar	Severity	Finding	Recommendation
--------	----------	---------	----------------

5. METHODOLOGY

VeldrixAI evaluates every AI response through a five-pillar trust framework. Each pillar applies a dedicated model or heuristic: Safety uses NVIDIA NemoGuard; Hallucination uses the HHEM-2.1 model; Bias uses distilroberta-bias; Prompt Security uses NemoGuard JailbreakDetect; and Compliance uses Microsoft Presidio for PII detection alongside regulatory rule sets. Pillar scores are weighted and combined into an overall trust score, which drives enforcement decisions.

6. RECOMMENDATIONS